



Moving Forward With Trust

Designing Human Accountability Into a Gen AI Workflow

Leadership Forum Community Summit 2026

Ken Judy · Senior Partner, Stride Consulting
kenjudy.us · stride.build



Artificial Intelligence ›

Effect on the Economy

A.I. and Campaigns

Transition for Workers

Anthropic's A.I. Model

China Closing the A.I. Gap

Xbox Hits Reset Button, Laying Off Thousands and Dropping Game Studios

The layoffs were part of wider cuts at Microsoft, as the company prioritizes spending on artificial intelligence.

The initial Xbox layoffs are part of **4,800 job cuts** that Microsoft announced on Monday, a total that accounts for roughly **2 percent of the company's work force**. It is Microsoft's latest employee culling as it plows tens of billions of dollars into the infrastructure for building artificial intelligence.

Software development is a design process



Manufacturing automates the repeated execution of a known solution. **D**esign is the activity of discovering the right solution under uncertainty.

When AI is applied to an undisciplined design process, it speeds up output without improving judgment. You get more of the wrong things faster.

(thank you Mary Poppendieck - author of Lean Software Development)

**Act on anything that
went wrong to avoid
errors of the same
nature in future**

**Plan what you are
doing**

Plan

Act

Do

Check

**Do what you said
you'd do**

**Check that you did it
right**

Developed by Walter A. Shewhart in the 1920's.
Popularized as part of the Toyota Production System.



"The Toyota style is not to create results by working hard. It is a system that says there is no limit to people's creativity. People don't go to Toyota to 'work' they go there to 'think'."

— Taichi Ohno

PDCA Cycle Applied to a Human In The Loop Agentic Workflow

Plan

Analyze the problem broadly tied to a positive user focused outcome

Task out the execution using templated guardrails

Agent is instructed to raise questions of clarification and present artifacts for approval.

Do

Execute in small atomic steps. Rely on programmatic deterministic steps where possible. Rely on Gen AI where fuzzy logic is helpful.

Enforce a feedback loop where the agent is showing its progress at each meaningful step.

Check

Completion analysis:
agent reviews session transcript and generated code against original plan. Verifies against an explicit definition of done.

Agent is instructed to report back with validation points not just a conclusion.

Act

Micro-retrospective.
Agent analyzes what worked and identifies targeted refinements to prompts and interaction patterns for the next cycle.

Agent is instructed to ask socratic questions to elicit insight from the human.

Step 1: Plan



What is the task, and what does a completed cycle deliver?

Example: *“Preparing and sending a weekly status update to your manager and stakeholders.”*

Step 2: Do



What requires human judgment that can't be delegated to AI?

Deciding how much detail on a sensitive risk to share with which audience — what goes to my manager versus the broader stakeholder version

What can AI execute autonomously, with review at defined gates?

AI drafts the metrics section from tracker data, drafts the risks list from open tickets, and drafts a first-pass narrative summary — human reviews all three before anything is sent

What are the irreversible or highest-stakes moments?

What AI behavior should trigger an immediate STOP?

Step 2: Do



What requires human judgment that can't be delegated to AI?

What can AI execute autonomously, with review at defined gates?

What are the irreversible or highest-stakes moments?

Sending the update to external stakeholders — once it's sent, I can't unsend a wrong number or a premature risk disclosure

What AI behavior should trigger an immediate STOP?

Stop if AI drafts a risk statement about a stakeholder's team without me reviewing the wording first. Stop if AI sends or schedules the update without my explicit sign-off. Stop if AI uses data older than this week without flagging it as stale.

Step 2: Do



What are the phase boundaries: "analysis done" and "execution done"?

Analysis done: this week's tracker data is pulled and I know which risks are still open. Execution done: the full draft exists with metrics, risks, and asks sections complete, before my own read-through.

Step 3: Check



Name 3-5 specific, observable signals of success.

(1) every shipped item this week is reflected accurately, (2) every open risk above medium severity is named with an owner, (3) every ask names a specific person and a deadline — zero vague asks, (4) I read and edited the final draft myself before it went out.

What are the most common failure modes to prevent?

A real risk gets buried in the details instead of called out up top. The update goes out without me actually reading the final draft. The numbers are stale because the draft used last week's data.

Step 4: Act



What process debt accumulates when this is repeated?

Skipping my own proofread pass when I'm busy. Reusing last week's risk list without checking whether anything actually resolved.

What do you want to learn and improve each cycle?

Track which flagged risks never actually mattered, so I get better at judging what's worth surfacing versus noise.



Access to AI workflow builder skill ▾

Share

(maybe from another conversation or a skill you were planning to build)?

use the pdca scaffold to guide me through creating a skill

Thinking about structuring skill development through Plan-Do-Check-Act cycles

Let's start with Layer 1.

Let's understand the task you want to systematize.

1. In one sentence: what is the task, and what does a completed cycle deliver? (For example: "Weekly client status report — delivers a reviewed, sent report with metrics and risks.")
2. How often is one full cycle performed? (daily / weekly / monthly / quarterly / on-demand / event-triggered)

Write a message...



Sonnet 5 Medium ▾



Plan

Think through
the questions

Do

Run the prompt
builder prompt

Check

Evaluate the
prompt by doing
test run(s)

Act

Use the agent to
help you refine
the prompt



Resources:

ken@kenjudy.us | kjudy@stride.build

Human directed ai workflow skill:

<https://github.com/kenjudy/human-directed-ai-workflow-builder>



AI Use Disclosure



I used Claude for brainstorming, argument review, and drafting assistance across multiple revisions. I personally verified all sources and made all final content decisions. I take full responsibility for the accuracy, originality, and quality of this work. [[References](#)]